


# ***Reinforcement learning***



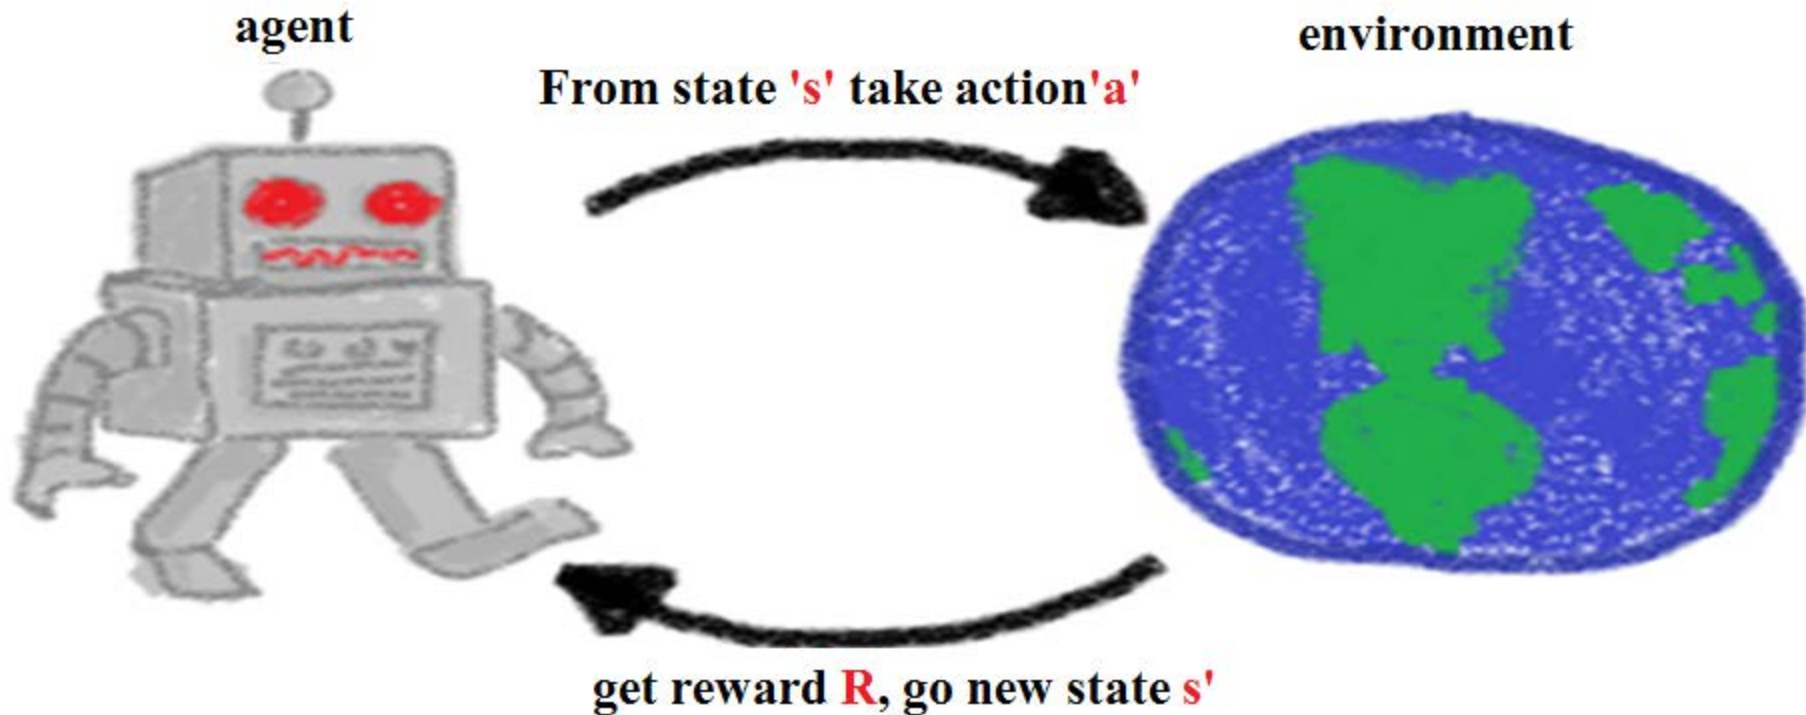
The term *reinforcement learning* comes from cognitive science and psychology and it describes the learning system of carrot and stick

Based on a motivational tactic that uses a reward and punishment system to encourage improved performance/behavior

learning by means of good or bad experience,

- **Reinforcement Learning (RL)**
- RL is learning through **direct experimentation**.
- It does not assume the existence of a teacher that provides **'training examples'** upon which learning of a task takes place
- Instead, in RL **experience** is the only teacher

# *Reinforcement learning*



Agent receives **no examples** and starts with **no model** of the environment.

Agent gets feedback through rewards, or **reinforcement**.

# Markov Decision Processes (MDP) describe an environment for reinforcement learning.

In MDPs the “**present state & action**” pair  
holds all the information required for  
determining **future state**

# Markov Decision Process (MDP)

- At each discrete time point
  - **Agent** observes state  $s_t$  and chooses **action**  $a_t$  (*according to some probability distribution*)
  - Receives **reward**  $r_t$  from the **environment** and the **state changes** to  $s_{t+1}$
- **Markov assumption:**  
$$r_t = r(s_t, a_t) \quad s_{t+1} = \delta(s_t, a_t)$$

$r_t$  and  $s_{t+1}$  depend only on the **current state and action**

A fundamental issue in RL algorithms is the balance/trade-off between **Exploration of the environment** and **Exploitation of information** already obtained by the agent.

### **Exploitation:**

Use current knowledge to take the optimal action

### **Exploration:**

Take possibly suboptimal action to gather more information about the environment

# Exploitation vs Exploration

**exploitation(greedy approach)** : Taking the action which the agent estimates to be the best at the current moment.

*the agent is exploiting its current knowledge about the reward structure of the environment to act.*

The opposite approach to greedy selection is to simply always take a **random action-exploration**.



- Restaurant Selection

- Exploitation: Go to favorite restaurant
- Exploration: Try a new restaurant

- Online Banner Advertisements

- Exploitation: Show the most successful advert
- Exploration: Show a different advert

- Clinical Trial

- Exploitation: Choose the best treatment so far
- Exploration: Try a new treatment

- Game Playing

- Exploitation: Play the move you believe is best
- Exploration: Play an experimental move

# Q-learning: learning the action-value function

Q-learning is about learning Q-values through observations.


**Q-learning** is a model-free reinforcement learning technique.

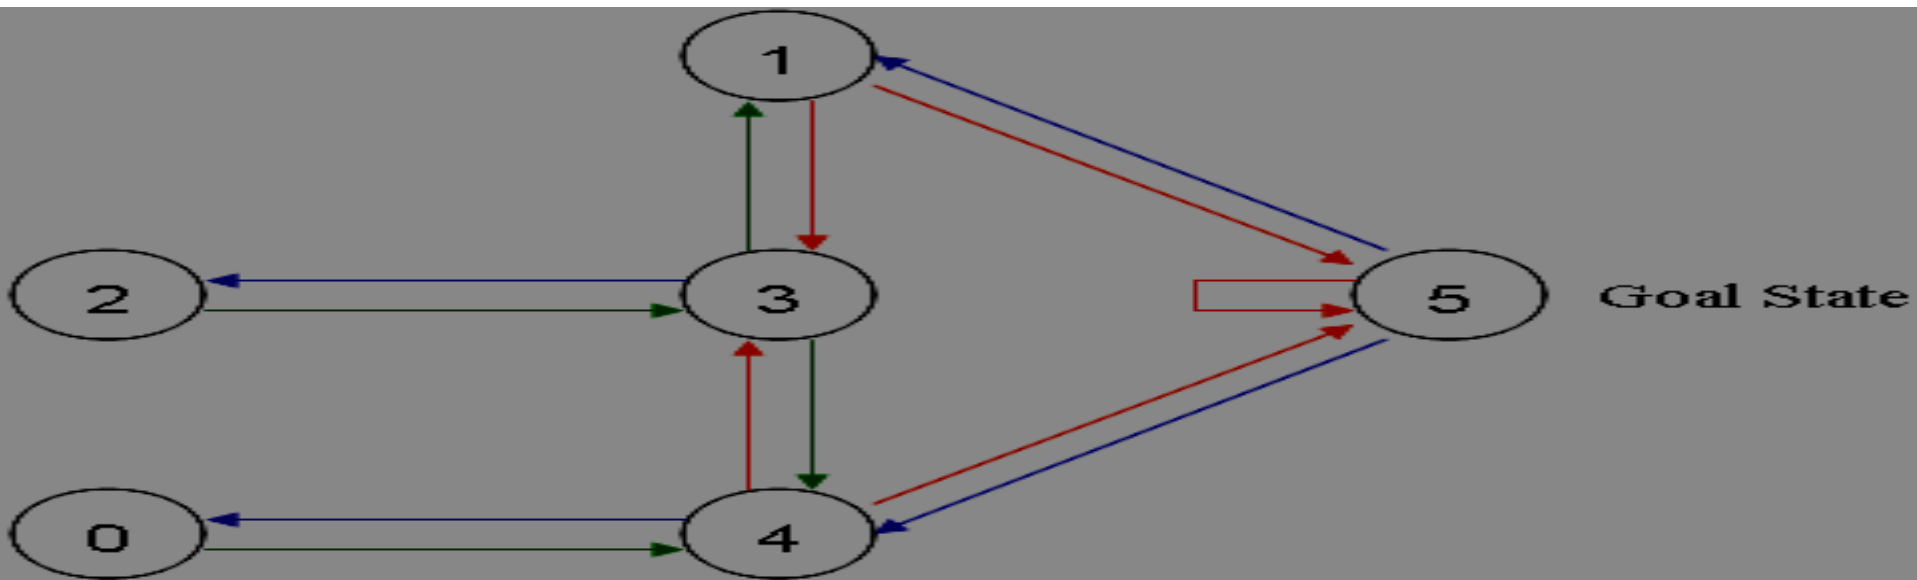

The 'Q' in Q-learning stands for quality.

Quality in this case represents how useful a given action is in gaining some future reward.

# Q Learning Example

- 5 rooms in a building connected by doors
- The outside of the building can be thought of as one big room (5).
- Doors 1 and 4 lead from building to room 5 (GOAL-outside).



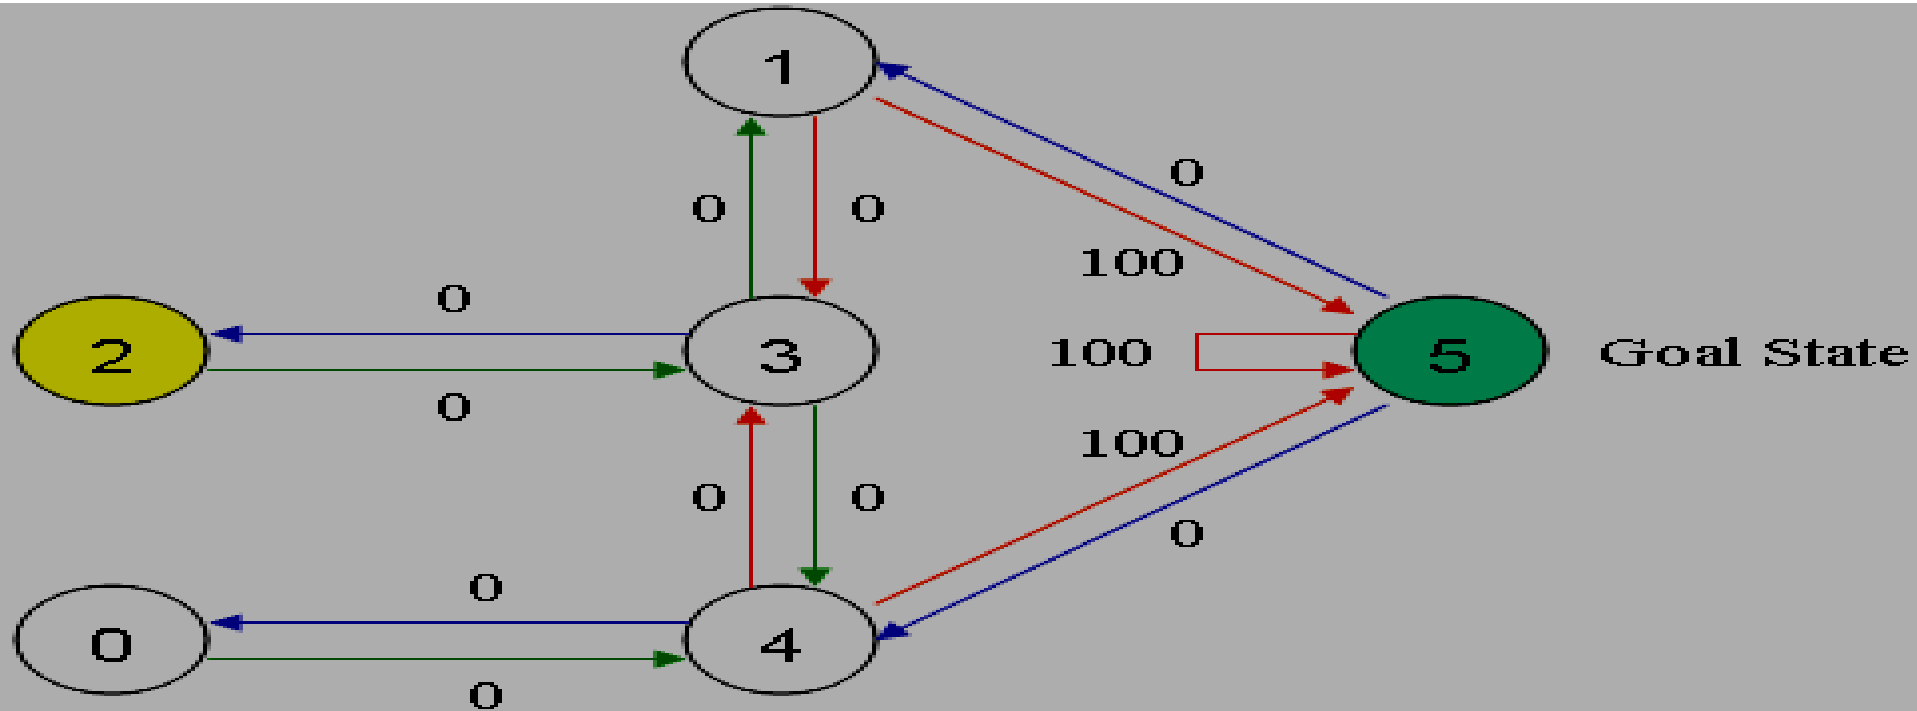


**Each room is a node**



**Each door is a link.**

- Goal room number =5
- Room 5 loops back to itself .

- Each Room, including outside (room no 5), is called a **"state"**
- The agent's movement from one room to another is an **"action"**




- In diagram, a **"state"** is depicted as a node, while **"action"** is represented by the arrows.




- Each arrow contains an instant reward value,
  - A reward value is associated to each door
  - The doors that lead immediately to the goal have an instant reward of 100
  - Otherwise reward = 0
- In Q-learning, **goal is to reach the state with the highest reward**, so that if the agent arrives at the goal, it will remain there forever, called an **"absorbing goal"**.

*•Imagine our agent as a dumb virtual robot that can learn through experience.*

**•The agent can pass from one room to another but has no knowledge of the environment, and doesn't know which sequence of doors lead to the outside.**



# Reward Table, "matrix R".



State

Action

0 1 2 3 4 5

0

1

2

3

4

5

$R =$

-1	-1	-1	-1	0	-1
-1	-1	-1	0	-1	100
-1	-1	-1	0	-1	-1
-1	0	0	-1	0	-1
0	-1	-1	0	-1	100
-1	0	-1	-1	0	100

**-1** : No link between nodes [ to distinguish between zero reward and no link]

**0** : reward is zero

**100** : reward is 100 if reaching to goal state

**MAKE Reward matrix**



## Matrix "Q"

- **Matrix, "Q"** : representing the **memory of agent** what the agent has learnt through experience.
- **Rows** : current state of the agent
- **Columns** : possible actions leading to the next state (the links between the nodes).
- Matrix Q is **initialized to zero** as:

$Q =$ 

present state	next states					
	0	1	2	3	4	5
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	0

$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \gamma * \text{Max}[Q(\text{next state}, \text{all actions})]$$

According to this formula, a value assigned to a specific element of matrix **Q**, is equal to the sum of the corresponding value in **matrix R** and the learning parameter **Gamma**, multiplied by the maximum value of **Q** for all possible actions in the next state.

$(0 \leq \gamma \leq 1)$  **closer to zero**  $\Rightarrow$  the agent will tend to consider only **immediate rewards**

**closer to one**  $\Rightarrow$  the agent will consider future rewards with greater weight, willing to **delay the reward**

# Algorithm

The Q-Learning algorithm goes as follows:

1. Set the gamma parameter, and environment rewards in matrix R.
2. Initialize matrix Q to zero
3. **For each episode:**

Select a random initial state.


**Do While** the goal state hasn't been reached.

- Select one among all possible actions for the current state.
- Using this possible action, consider going to the next state.
- Get maximum Q value for this next state based on all possible actions.
- Compute:  $Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$
- Set the next state as the current state.

**End Do**

**End For**



- Set learning parameter  $\text{Gamma} = 0.8$ ,
- **initial state as Room 1.**
- Initialize matrix  $Q$  as a zero matrix:



$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

$$R = \begin{matrix} & \text{Action} \\ \text{State} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix} \end{matrix}$$

**By random selection, select to go to 5 as action**



$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$$

$$Q(1, 5) = R(1, 5) + 0.8 * \text{Max}[Q(5, 1), Q(5, 4), Q(5, 5)] = 100 + 0.8 * 0 = 100$$

- Since matrix  $Q$  is still initialized to zero,  $Q(5, 1)$ ,  $Q(5, 4)$ ,  $Q(5, 5)$ , are all zero.
- The next state, 5, now becomes the current state.
- **Because 5 is the goal state, finished one episode.**
- agent's brain now contains an updated matrix  $Q$  as:

$$Q = \begin{matrix} & 0 & 1 & 2 & 3 & 4 & 5 \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

$$Q = \begin{matrix} & 0 & 1 & 2 & 3 & 4 & 5 \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

- For **next episode**, randomly chose initial state.
- *choose state 3 as initial state.*
- By random selection, *select state 1 as action.*



$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$

$$Q(3, 1) = R(3, 1) + 0.8 * \text{Max}[Q(1, 3), Q(1, 5)]$$

$$= 0 + 0.8 * \text{Max}(0, 100) = 80$$

update matrix Q

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$




The matrix Q becomes

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 80 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

- The next **state, 1, now becomes the current state.**
- Repeat the inner loop of the Q learning algorithm **because state 1 is not the goal state[ episode not over].**



- Set current state as 1
- two possible actions: go to state 3, or go to state 5.
- Randomly **choose action state 5.**





$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$$

$$Q(1, 5) = R(1, 5) + 0.8 * \text{Max}[Q(5, 1), Q(5, 4), Q(5, 5)] = 100 + 0.8 * 0 = 100$$



- Reached to goal state **finish this episode.**
- agent's brain now contain updated matrix Q as:

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 80 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

- Agent learns more through further episodes
- finally reach convergence values in matrix Q like:

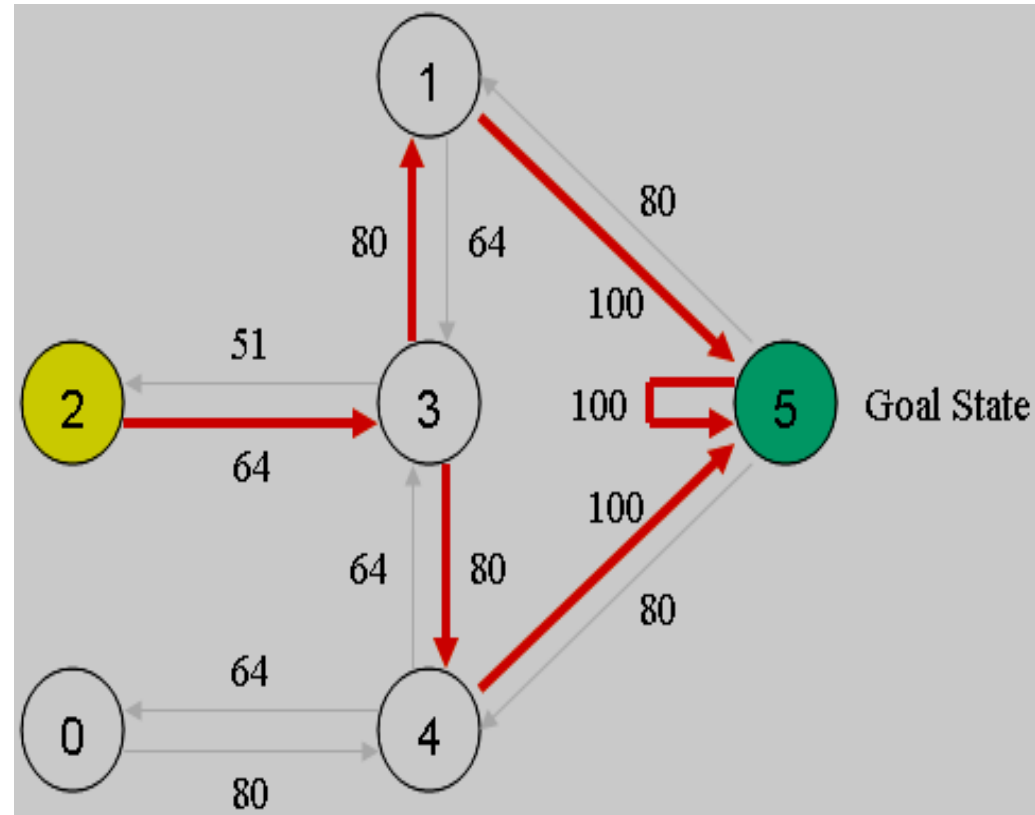
$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 400 & 0 \\ 0 & 0 & 0 & 320 & 0 & 500 \\ 0 & 0 & 0 & 320 & 0 & 0 \\ 0 & 400 & 256 & 0 & 400 & 0 \\ 320 & 0 & 0 & 320 & 0 & 500 \\ 0 & 400 & 0 & 0 & 400 & 500 \end{bmatrix} \end{matrix}$$

This matrix Q, can then be **normalized** (i.e.; converted to percentage) by dividing all non-zero entries by 5 :

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 80 & 0 \\ 0 & 0 & 0 & 64 & 0 & 100 \\ 0 & 0 & 0 & 64 & 0 & 0 \\ 0 & 80 & 51 & 0 & 80 & 0 \\ 64 & 0 & 0 & 64 & 0 & 100 \\ 0 & 80 & 0 & 0 & 80 & 100 \end{bmatrix} \end{matrix}$$

- Agent has learned the most optimal paths to the goal state through experience.
- Best sequences : the links with the **highest values** at each state.

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 80 & 0 \\ 0 & 0 & 0 & 64 & 0 & 100 \\ 0 & 0 & 0 & 64 & 0 & 0 \\ 0 & 80 & 51 & 0 & 80 & 0 \\ 64 & 0 & 0 & 64 & 0 & 100 \\ 0 & 80 & 0 & 0 & 80 & 100 \end{bmatrix} \end{matrix}$$



If we start at 2,  
the best sequence is 2 - 3 - 1 – 5 or 2-3-4-5

# DEEP Q LEARNING

