

# **Process Synchronization**

# Background

- Cooperation is required for several reasons like:  
Information sharing, Computation speedup, Modularity,  
Convenience
- Concurrent access to shared data may result in data inconsistency.
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.

# Bounded-buffer

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    ...
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
int counter = 0;
```

# Bounded-buffer

- **Producer process**

**item nextProduced;**

**while (1) {**

**while (counter == BUFFER\_SIZE)**

**; /\* do nothing \*/**

**buffer[in] = nextProduced;**

**in = (in + 1) % BUFFER\_SIZE;**

**counter++;**

**}**

# Bounded-buffer

## ■ Consumer process

```
item nextConsumed;
```

```
while (1) {  
    while (counter == 0)  
        ; /* do nothing */  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
}
```

# Bounded Buffer

- The statements

**counter++;**  
**counter--;**

must be performed *atomically*.

- Atomic operation means an operation that completes in its entirety without interruption.



# Bounded Buffer

- The statement “**count++**” may be implemented in machine language as:

**register1 = counter**

**register1 = register1 + 1**

**counter = register1**

- The statement “**count--**” may be implemented as:

**register2 = counter**

**register2 = register2 – 1**

**counter = register2**

# Bounded Buffer

- If both the producer and consumer attempt to update the buffer concurrently, the assembly language statements may get interleaved.
- Interleaving depends upon how the producer and consumer processes are scheduled.



# Bounded Buffer

- Assume **counter** is initially 5. One interleaving of statements is:

**producer:** **register1 = counter** (*register1 = 5*)

**producer:** **register1 = register1 + 1** (*register1 = 6*)

**consumer:** **register2 = counter** (*register2 = 5*)

**consumer:** **register2 = register2 - 1** (*register2 = 4*)

**producer:** **counter = register1** (*counter = 6*)

**consumer:** **counter = register2** (*counter = 4*)

- The value of **count** may be either 4 or 6, where the correct result should be 5.

# Race Condition

- **Race condition:** The situation where several processes access – and manipulate shared data concurrently. The final value of the shared data depends upon which process finishes last.
- To prevent race conditions, concurrent processes must be **synchronized**.

# The Critical-section Problem

- $n$  processes all competing to use some shared data
- Each process has a code segment, called *critical section*, in which the shared data is accessed.
- **Problem** – ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section.

# Solution to Critical-section Problem

1. **Mutual Exclusion**: If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections.
2. **Progress**: If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely. (~ Deadlock)
3. **Bounded Waiting**: A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted. (~ Starvation)



# Requirements for Solution

- Only one process at a time is allowed in the critical section for a resource
- A process that halts in its non-critical section must do so without interfering with other processes
- No deadlock or starvation
- A process must not be delayed access to a critical section when there is no other process using it
- No assumptions are made about relative process speeds or number of processes
- A process remains inside its critical section for a finite time only

# Initial Attempts to Solve Problem

- Only 2 processes,  $P_0$  and  $P_1$
- General structure of process  $P_i$  (other process  $P_j$ )

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (1);
```

- Processes may share some common variables to synchronize their actions.



# Algorithm 1

- Shared variables:

- **int turn;**

- initially **turn = 0**

- **turn = i**  $\Rightarrow P_i$  can enter its critical section

- Process  $P_i$

- do {**

- while (turn != i) do nothing;**

- critical section

- turn = j;**

- remainder section

- } while (1);**

- Satisfies mutual exclusion, but not progress

- This solution guarantees mutual exclusion
- Drawback 1: processes must strictly alternate
- Drawback 2: if one processes fails other process is permanently blocked
- This problem arises due to fact that it stores name of the process that may enter critical section rather than the process state

# Algorithm 2

- Shared variables

- **boolean flag[2];**  
initially **flag [0] = flag [1] = false.**
- **flag [i] = true**  $\Rightarrow P_i$  ready to enter its critical section

- Process  $P_i$

do {

**flag[i] := true;**  
**while (flag[j]) do nothing;**

critical section

**flag [i] = false;**

remainder section

} while (1);

- Satisfies mutual exclusion, but not progress requirement.

- This approach Satisfies mutual exclusion
- This approach may lead to dead lock

What is wrong with this implementation ?

- A process sets its state without knowing the state of other. Dead lock occurs because each process can insist on its right to enter critical section

# Algorithm 3 (Peterson's Solution)

- Combined shared variables of algorithms 1 and 2.
- **Process  $P_i$** 
  - do {**
    - flag [i] := true;**
    - turn = j;**
    - while (flag [j] and turn = j) do nothing;**
    - critical section
    - flag [i] = false;**
    - remainder section
  - } while (1);**
- Meets all three requirements; solves the critical-section problem for two processes.



# Bakery Algorithm

## Critical section for n processes

- Before entering its critical section, process receives a number. Holder of the smallest number enters the critical section.
- If processes  $P_i$  and  $P_j$  receive the same number, if  $i < j$ , then  $P_i$  is served first; else  $P_j$  is served first.
- The numbering scheme always generates numbers in increasing order of enumeration; i.e., 1,2,3,3,3,3,4,5...



# Bakery Algorithm

- Notation  $\leq$  lexicographical order (ticket #, process id #)
  - $(a,b) < (c,d)$  if  $a < c$  or if  $a = c$  and  $b < d$
  - $\max(a_0, \dots, a_{n-1})$  is a number,  $k$ , such that  $k \geq a_i$  for  $i = 0, \dots, n - 1$

- Shared data

**boolean choosing[n];**

**int number[n];**

Data structures are initialized to **false** and **0** respectively

# Bakery Algorithm

do {

```
    choosing[i] = true;
```

```
    number[i] = max(number[0], number[1], ..., number [n – 1])+1;
```

```
    choosing[i] = false;
```

```
    for (j = 0; j < n; j++) {
```

```
        while (choosing[j]) do no-op ;
```

```
        while ((number[j] != 0) && ((number[j],j) < (number[i],i)))
```

```
            do no-op ;
```

```
    }
```

critical section

```
    number[i] = 0;
```

remainder section

```
} while (1);
```

# Mutual Exclusion: Hardware Support

## ■ Interrupt Disabling

- ➔ A process runs until it invokes an operating-system service or until it is interrupted
- ➔ Disabling interrupts guarantees mutual exclusion
- ➔ Processor is limited in its ability to interleave programs
- ➔ Multiprocessing
  - disabling interrupts on one processor will not guarantee mutual exclusion

# Mutual Exclusion: Hardware Support

- Special Machine Instructions
  - ➔ Performed in a single instruction cycle (Atomic)
  - ➔ Not subject to interference from other instructions
  - ➔ Test and Set
  - ➔ Swap

# Synchronization Hardware

- Test and modify the content of a word atomically

```
boolean TestAndSet(boolean &target)
{
    boolean rv = target;
    target = true;
    return rv;
}
```

# Mutual Exclusion With Test-and-set

- Shared data:

**boolean lock = false;**

- Process  $P_i$

**do {**

**while (TestAndSet(lock)) do no-op;**

critical section

**lock = false;**

remainder section

**}**



# Bounded-waiting Mutual Exclusion with TestAndSet()

do {

```
waiting[i] = TRUE;  
key = TRUE;  
while (waiting[i] && key)  
    key = TestAndSet(&lock);  
waiting[i] = FALSE;
```

// critical section

```
j = (i + 1) % n;  
while ((j != i) && !waiting[j])  
    j = (j + 1) % n;  
if (j == i)  
    lock = FALSE;  
else  
    waiting[j] = FALSE;
```

// remainder section

} while (TRUE);

# Synchronization Hardware

- Atomically swap two variables.

```
void Swap(boolean &a, boolean &b)
{
    boolean temp = a;
    a = b;
    b = temp;
}
```

# Mutual Exclusion With Swap

- Shared data (initialized to **false**):  
**boolean lock;**

- Process  $P_i$

do {

```
key = true;  
while (key == true)  
    Swap(lock, key);
```

critical section

```
lock = false;
```

remainder section

}

- Key is a local boolean variable

# Mutual Exclusion Machine Instructions

## ■ Advantages

- ➔ Applicable to any number of processes on either a single processor or multiple processors. It is simple and therefore easy to verify
- ➔ It can be used to support multiple critical sections

# Mutual Exclusion Machine Instructions

- Disadvantages

- ➔ Busy-waiting consumes processor time
- ➔ Starvation is possible when a process leaves a critical section and more than one process is waiting.

# Semaphores (OS Support)

- Semaphore is a variable that has an integer value
  - ➔ May be initialized to a nonnegative number
  - ➔ Wait operation decrements the semaphore value
  - ➔ Signal operation increments semaphore value
- Wait and signal operations cannot be interrupted
- If a process is waiting for a signal, it is suspended until that signal is sent
- Queue is used to hold processes waiting on the semaphore



# Semaphores

- Synchronization tool that **does not require busy waiting**.
- Semaphore  $S$  – integer variable
- can only be accessed via two indivisible (atomic) operations

*wait* ( $S$ ):

**while  $S \leq 0$  do *no-op*;**  
 **$S--$ ;**

---

*signal* ( $S$ ):

**$S++$ ;**

- ***In the above definition the value of semaphore is never negative***

# Critical Section of $N$ Processes

- Shared data:

**semaphore mutex;**    *//initially mutex = 1*

- Process  $P_i$ :

do {

**wait(mutex);**

    critical section

**signal(mutex);**

    remainder section

} while (1);

- **SPINLOCKS (Busy Waiting):** *Useful only if waiting period is small and hence may save a context switch*

# **Avoiding Busy Waiting**

# Semaphore Implementation

- Define a semaphore as a record

```
typedef struct {  
    int value;  
    struct process *L;  
} semaphore;
```

- Assume two simple operations:

→ **block** suspends the process that invokes it.

→ **wakeup(*P*)** resumes the execution of a blocked process *P*.

# Implementation

- Semaphore operations are now defined as

*wait(S):*

```
S.value--;  
if (S.value < 0) {  
    add this process to S.L;  
    block;  
}
```

---

*signal(S):*

```
S.value++;  
if (S.value <= 0) {  
    remove a process P from S.L;  
    wakeup(P);  
}
```

- Value of semaphore can be negative and represents the number of processes waiting on it



# Semaphore As a General Synchronization Tool

- Execute  $B$  in  $P_j$  only after  $A$  executed in  $P_i$
- Use semaphore  $flag$  initialized to 0
- Code:

$P_i$	$P_j$
$\vdots$	$\vdots$
$A$	$wait(flag)$
$signal(flag)$	$B$

- Semaphore operations are now defined as

```
wait(S):
    S.value--;
    if (S.value < 0) {
        add this process to S.L;
        block;
    }

signal(S):
    S.value++;
    if (S.value <= 0) {
        remove a process P from S.L;
        wakeup(P);
    }
```

- Value of semaphore can be negative and represents the number of processes waiting on it

# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.
- Let S and Q be two semaphores initialized to 1

$P_0$	$P_1$
<i>wait(S);</i>	<i>wait(Q);</i>
<i>wait(Q);</i>	<i>wait(S);</i>
$\vdots$	$\vdots$
<i>signal(S);</i>	<i>signal(Q);</i>
<i>signal(Q)</i>	<i>signal(S);</i>

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

# Two Types of Semaphores

- *Counting* semaphore – integer value can range over an unrestricted domain.
- *Binary* semaphore – integer value can range only between 0 and 1; can be simpler to implement.
- Can implement a counting semaphore  $S$  as a binary semaphore.

# Classical Problems of Synchronization

- Bounded-Buffer Problem
- Dining-Philosophers Problem
- Readers-Writers Problem

# Bounded-buffer Problem

- Shared data

**semaphore full, empty, mutex;**

Initially:

**full = 0, empty = n, mutex = 1**

- Buffer size is n
- Mutex provides exclusive access to the buffer
- Consumers wait on full
- Producers wait on empty




# Bounded-buffer Problem: Producer Process

```
do {  
    ...  
    produce an item in nextp  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    add nextp to buffer  
    ...  
    signal(mutex);  
    signal(full);  
} while (1);
```

# Bounded-buffer Problem: Consumer Process

```
do {  
    wait(full)  
    wait(mutex);  
    ...  
    remove an item from buffer to nextc  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    consume the item in nextc  
    ...  
} while (1);
```

# Dining-philosophers Problem



- Shared data

**`semaphore chopstick[5];`**

Initially all values are 1

- **To start eating, a philosopher needs two chopsticks**
- **A philosopher may pick up only one chopstick at a time**
- **After eating, the philosopher releases both the chopsticks**

# Dining-philosophers Problem

■ Philosopher  $i$ :

do {

```
wait(chopstick[i])  
wait(chopstick[(i+1) % 5])
```

...

eat

...

```
signal(chopstick[i]);  
signal(chopstick[(i+1) % 5]);
```

...

think

...

} while (1);



**The solution is not deadlock free!!**

**→ All philosophers pick up left chopsticks!!**

---

**→ Allow at most 4 philosophers to be sitting on the table**

**→ Allow a philosopher to pick chopsticks only if both are available**

**→ An *odd* philosopher picks up first the left and then the right chopstick while a *even* philosopher does the reverse**

# Implementation

- Semaphore operations are now defined as

*wait(S):*

```
S.value--;  
if (S.value < 0) {  
    add this process to S.L;  
    block;  
}
```

---

*signal(S):*

```
S.value++;  
if (S.value <= 0) {  
    remove a process P from S.L;  
    wakeup(P);  
}
```

- Value of semaphore can be negative and represents the number of processes waiting on it

# Readers-Writers Problem

- Two readers can access the shared data item simultaneously.
- A writer requires exclusive access
- No reader should wait unless a writer is already in critical section
- Writers may starve
- **Shared** data

```
var mutex, wrt: semaphore (=1);  
    readcount : integer (=0);
```

- **wrt** is common to both readers and writers. It functions as mutual exclusion semaphore for writers. It is also used by first and last reader that enters or exits the CS.
- **readcount** keeps track of how many readers are currently accessing the object.
- **mutex** provides mutual exclusion for updating *readcount*.

# Readers-Writers Problem

## ■ Writer process

*wait(wrt);*

...

writing is performed

...

*signal(wrt);*

# Readers-Writers Problem (Cont.)

## ■ Reader process

```
wait(mutex);  
    readcount := readcount + 1;  
    if readcount = 1 then wait(wrt);  
signal(mutex);  
    ...  
    reading is performed  
    ...  
wait(mutex);  
    readcount := readcount - 1;  
    if readcount = 0 then signal(wrt);  
signal(mutex);
```



# Monitors


- Monitor is a software module
- Chief characteristics
  - ➔ Local data variables are accessible only by the monitor
  - ➔ Process enters monitor by invoking one of its procedures
  - ➔ Only one process may be executing in the monitor at a time

# Monitors

- High-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes.

```
type monitor-name = monitor  
    variable declarations  
    procedure entry  $P1$  : (...);  
        begin ... end;  
    procedure entry  $P2$  (...);  
        begin ... end;  
        ⋮  
    procedure entry  $Pn$  (...);  
        begin...end;  
begin  
    initialization code  
end
```

# Schematic view of a monitor



# Monitors (Cont.)

- To allow a process to wait within the monitor, a *condition* variable must be declared, as

**var** *x, y: condition*

- Condition variable can only be used with the operations *wait* and *signal*.

→ The operation


*x.wait;*

means that the process invoking this operation is suspended until another process invokes


*x.signal;*

→ The *x.signal* operation resumes exactly **one** suspended process. If no process is suspended, then the signal operation has no effect.

# Monitor with condition variables







# Dining Philosophers Example

**type** *dining-philosophers* = **monitor**

**var** *state* : **array** [0..4] **of** :(*thinking*, *hungry*, *eating*);

**var** *self* : **array** [0..4] **of** *condition*;

**procedure** *pickup* (*i*: 0..4);

**begin**

*state*[*i*] := *hungry*;

*test* (*i*);

**if** *state*[*i*] ≠ *eating* **then** *self*[*i*].*wait*;

**end**;

**procedure** *putdown* (*i*: 0..4);

**begin**

*state*[*i*] := *thinking*;

*test* (*i*+4 **mod** 5);

*test* (*i*+1 **mod** 5);

**end**;

# Dining Philosophers (Cont.)

```
procedure test(i: 0..4);  
  begin  
    if state[i+4 mod 5]  $\neq$  eating  
      and state[i] = hungry  
      and state[i+1 mod 5]  $\neq$  eating  
    then begin  
      state[i] := eating;  
      self[i].signal;  
    end;  
  
  end;
```

```
  begin  
    for i := 0 to 4  
      do state[i] := thinking;  
    end.
```

**Philosopher  $i$  must invoke the operations *pickup* and *putdown* in the following sequence:**

**dp.pickup(i)**

**.....**

**Eat**

**.....**

**dp.putdown(i)**

# Monitor Implementation Using Semaphores

## ■ Variables

```
var mutex: semaphore (init = 1)  
    next: semaphore (init = 0)  
    next-count: integer (init = 0)
```

## ■ Each external procedure $F$ will be replaced by $wait(mutex);$

```
    ...  
    body of  $F$ ;
```

```
    ...  
if next-count > 0  
    then signal(next)  
    else signal(mutex);
```

## ■ Mutual exclusion within a monitor is **ensured**.

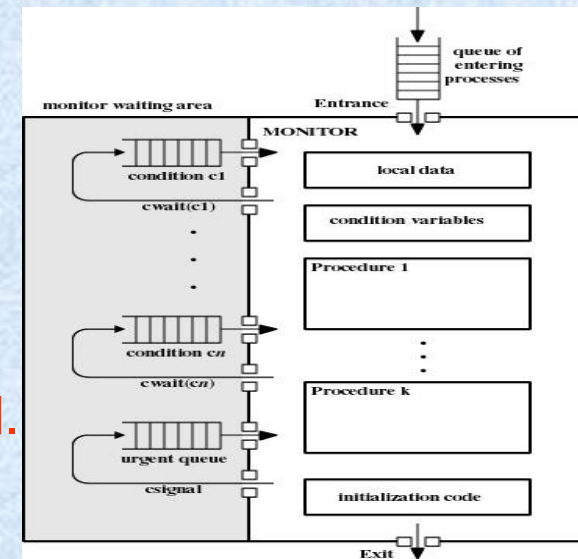


Figure 5.21 Structure of a Monitor



# Monitor Implementation (Cont.)

- For each condition variable  $x$ , we have:

**var**  $x$ -sem: semaphore (init = 0)

$x$ -count: integer (init = 0)

- The operation  $x$ .wait can be implemented as:

$x$ -count :=  $x$ -count + 1;

**if** next-count > 0

**then** signal(next)

**else** signal(mutex);

wait( $x$ -sem);

$x$ -count :=  $x$ -count - 1;




Figure 5.21 Structure of a Monitor

# Monitor Implementation (Cont.)

- The operation `x.signal` can be implemented as:

**if** `x-count` > 0

**then begin**

`next-count := next-count + 1;`

`signal(x-sem);`

`wait(next);`

`next-count := next-count - 1;`

**end;**

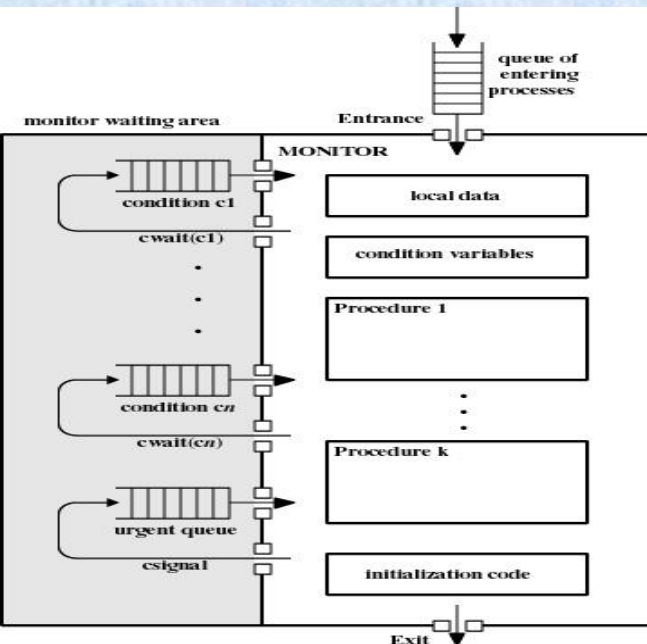


Figure 5.21 Structure of a Monitor

# Monitor Implementation (Cont.)

- *Conditional-wait* construct:  $x.wait(c);$ 
  - $c$  – integer expression evaluated when the wait operation is executed.
  - value of  $c$  (*priority number*) stored with the name of the process that is suspended.
  - when  $x.signal$  is executed, process with **smallest** associated priority number is resumed next.